

The emergence of words from vocal imitations

Journal:	<i>Proceedings B</i>
Manuscript ID	RSPB-2017-1610
Article Type:	Research
Date Submitted by the Author:	18-Jul-2017
Complete List of Authors:	Edmiston, Pierce; University of Wisconsin-Madison, Department of Psychology Perlman, Marcus; University of Wisconsin-Madison, Psychology Lupyan, Gary; University of Wisconsin Madison, Psychology
Subject:	Cognition < BIOLOGY, Evolution < BIOLOGY, Behaviour < BIOLOGY
Keywords:	language evolution, iconicity, vocal imitation, transmission chain
Proceedings B category:	Behaviour

SCHOLARONE™
Manuscripts

1

2

3

4

5

6

9

10

Abstract

People have long pondered the origins of language, especially the words that compose them. Here, we report a series of experiments investigating how conventional spoken words might emerge from imitations of environmental sounds. Does the repeated imitation of an environmental sound gradually give rise to novel word forms? In what ways do these words resemble the original sounds that motivated them? Participants played a version of the children's game "Telephone". The first generation of participants imitated recognizable environmental sounds (e.g., glass breaking, water splashing). Subsequent generations imitated the imitations for a maximum of 8 generations. The results showed that the imitations became more stable and word-like, and later imitations were easier to learn as category labels. At the same time, even after 8 generations, both spoken imitations and their written transcriptions could be matched above chance to the category of environmental sound that motivated them. These results show how repeated imitation can create progressively more word-like forms while continuing to retain a resemblance to the original sound that motivated them, and speak to the possible role of human vocal imitation in explaining the origins of at least some spoken words.

Keywords: language evolution, iconicity, vocal imitation, transmission chain

Word count: 5552

The importance of imitation and depiction in the origin of signs is clearly observable in signed languages (Goldin-Meadow, 2016; Kendon, 2014; Klima & Bellugi, 1980), but in considering the idea that imitation in the vocal modality may be key to understanding the origin of spoken words, many have argued that the human capacity for vocal imitation is far too limited to play a significant role (Arbib, 2012; Armstrong & Wilcox, 2007; Corballis, 2003; Hewes, 1973; Hockett, 1978; Tomasello, 2010). For example, Pinker and Jackendoff (2005) argued that, “most humans lack the ability... to convincingly reproduce environmental sounds... Thus ‘capacity for vocal imitation’ in humans might be better described as a capacity to learn to produce speech” (p. 209). Consequently, it is still widely assumed that vocal imitation — or more broadly, the use of any sort of resemblance between form and meaning — cannot be important to understanding the origin of spoken words. We challenge this view by demonstrating that spoken words can emerge from vocal imitations even without the intention to communicate. We find that repeating vocal imitations of environmental sounds over generations of unique speakers is sufficient to create more word-like vocalizations both in form and function.

Although most words of contemporary spoken languages are not clearly imitative in origin, there has been a growing recognition of the importance of imitative words in spoken languages (Dingemanse, Blasi, Lupyan, Christiansen, & Monaghan, 2015; Perniss, Thompson, & Vigliocco, 2010) and the frequent use of vocal imitation and depiction in spoken discourse (Clark & Gerrig, 1990; Lewis, 2009). This has led some to argue for the importance of imitation for understanding the origin of spoken words (e.g., Brown, Black, & Horowitz, 1955; Dingemanse, 2014; Donald, 2016; Imai & Kita, 2014; Perlman, Dale, & Lupyan, 2015). In addition, counter to previous assumptions, people are highly effective at using vocal imitations to refer to environmental sounds such as coins dropping in a jar or mechanical events such as scraping — in some cases, even more effective than when using conventional words (Lemaitre & Rocchesso, 2014). Recent work has also shown that people

WORDS FROM IMITATIONS

4

are able to create novel imitative vocalizations for more abstract meanings (e.g. “slow”, “rough”, “good”, “many”) that are understandable to naïve listeners (Perlman et al., 2015). These imitations are effective not because people can mimic environmental sounds with high fidelity, but because people are able to produce imitations that capture the salient features of sounds in ways that are understandable to listeners (Lemaitre, Houix, Voisin, Misdariis, & Susini, 2016). Similarly, the features of onomatopoeic words might highlight distinctive aspects of the sounds they represent. For example, the initial voiced, plosive /b/ in “boom” represents an abrupt, loud onset, the back vowel /u/ a low pitch, and the nasalized /m/ a slow, muffled decay (Rhodes, 1994).

Thus, converging evidence suggests that people can use vocal imitation as an effective means of communication. But can vocal imitations ever give rise to words that can be integrated into the vocabulary of a language? And if so, by what means might this happen? To answer these questions, we recruited participants to play an online version of the children’s game of “Telephone”. In the children’s game, a spoken message is whispered from one person to the next. In our version, the original message or “seed sound” was a recording of an environmental sound. The initial group of participants (first generation) imitated these seed sounds, the next generation imitated the previous imitators, and so on for up to 8 generations.

We then conducted a series of analyses and additional experiments to systematically answer the following questions: First, do imitations stabilize in form and become more word-like as they are repeated? Second, do the imitations retain a resemblance to the original environmental sound that inspired them? If so, it should be possible for naïve participants to match the emergent words back to the original seed sounds. Third, do the imitations become more suitable as labels for the category of sounds that motivated them? For example, does the imitation of a particular water-splashing sound become, over generations of repeated imitation, a better label for the more general category of water-splashing sounds?

Experiment 1: Stabilization of imitations through repetition

In the first experiment, we collected the vocal imitations, and assessed the extent to which repeating imitations of environmental sounds over generations of unique speakers results in progressive stabilization toward more word-like forms. After collecting the imitations, we measured changes in the stability of the imitations in three ways. First, we measured changes in the perception of acoustic similarity between subsequent generations of imitations along contiguous transmission chains. Second, we used algorithmic measures of acoustic similarity to assess the similarity of imitations sampled within and between transmission chains. Third, we obtained transcriptions of imitations, and measured the extent to which later generation imitations were transcribed with greater consistency and agreement. The results show that repeated imitation results in vocalizations that are easier to repeat with high fidelity and easier to transcribe into English orthography.

Methods

Selecting seed sounds. To avoid sounds having lexicalized or conventionalized onomatopoeic forms in English, we used inanimate categories of environmental sounds. Using an odd-one-out norming procedure ($N=105$ participants), an initial set of 36 sounds in 6 categories was reduced to a final set of 16 “seed” sounds: 4 sounds in each of 4 categories. The purpose of this norming procedure was to reach a set of approximately equally distinguishable sounds within each category by systematically removing the sounds that stood out in each category. The results of the norming procedure are shown in Fig. S1. The four final categories were: water, glass, tear, zipper. The final 16 seed sounds can be downloaded from here: osf.io/n6g7d/download.

Collecting vocal imitations. Participants ($N=94$) recruited from Amazon Mechanical Turk were paid to participate in an online version of the children’s game of “Telephone”. Participants were instructed that they would hear some sound and their task is to reproduce it as accurately as possible using their computer microphone. Full instructions

are provided in the Supplemental Materials.

Each participant listened to and imitated four sounds: one from each of the four categories of environmental sounds. Sounds were assigned at random such that participants were unlikely to imitate the same person more than once. Participants were allowed to listen to each target sound multiple times, but were only allowed a single recording in response. Recordings that were too quiet (less than -30 dBFS) were not accepted.

Imitations were monitored by an experimenter to catch any gross errors in recording before they were heard by the next generation of imitators. For example, recordings with loud sounds in the background were removed, and recordings were trimmed to the length of the imitation prior to the next generation. The experimenter also removed sounds that violated the rules of the experiment, e.g., by saying something in English. A total of 115 (24%) imitations were removed prior to subsequent analysis. The final sample contained 365 imitations along 105 contiguous transmission chains (Fig. 1).

Measuring acoustic similarity.

Acoustic similarity judgments. Acoustic similarity judgments were gathered from five research assistants who listened to pairs of sounds (approx. 300) and rated their subjective similarity. On each trial, raters heard two sounds from subsequent generations played in random order. They then indicated the similarity between the sounds on a 7-point Likert scale from *Entirely different and would never be confused* to *Nearly identical*. Raters were encouraged to use as much of the scale as they could while maximizing the likelihood that, if they did this procedure again, they would reach the same judgments. Full instructions are provided in the Supplemental Materials. Inter-rater reliability was calculated as the intra-class coefficient treating the group as the unit of analysis (Gamer, Lemon, Fellows, & Singh, 2012; Shrout & Fleiss, 1979): $ICC = 0.76$, 95% CI [0.70, 0.81], $F(170, 680) = 4.18$, $p < 0.001$. Ratings were normalized for each rater (z-scored) prior to analysis.

Algorithmic acoustic similarity. To obtain algorithmic measures of acoustic similarity, we used the acoustic distance functions included in Phonological Corpus Tools




Figure 1. Vocal imitations collected in the transmission chain experiment. Seed sounds (16) were sampled from four categories of environmental sounds: glass, tear, water, zipper. Participants imitated each seed sound, and then the next generation of participants imitated the imitations, and so on, for up to 8 generations. Chains are unbalanced due to random assignment and the exclusion of some low quality recordings.

(Hall, Allen, Fry, Mackie, & McAuliffe, 2016). We computed Mel-frequency cepstral coefficients (MFCCs) between pairs of imitations using 12 coefficients in order to obtain speaker-independent estimates.

Collecting transcriptions of imitations. Participants ($N=216$) recruited from Amazon Mechanical Turk were paid to transcribe vocalizations using English orthography, being instructed to write down what they heard as a single “word” so that the written word would sound as much like the sound as possible. Participants were instructed that this was a word creation task and so to avoid transcribing the vocalizations into existing English words. Each participant completed 10 transcriptions. Transcriptions were gathered for the first and the last three generations of imitations collected in the transmission chain experiment. Participants also provided “transcriptions” of the original environmental seed sounds. Analyses of these transcriptions are reported in the Supplementary Materials (Fig. S5).

To measure similarity among transcriptions of the same imitation, we used the `SequenceMatcher` functions in the `difflib` package of the Python standard library, which implements a version of Ratcliff and Obershelp’s “gestalt pattern matching” algorithm. Alternative measures of transcription agreement including exact string matching and the length of the longest substring match were also collected.

Analyses. Statistical analyses were conducted in R using linear mixed-effects models provided by the `lme4` package (Bates, Mächler, Bolker, & Walker, 2015). Degrees of freedom and corresponding significance tests for linear mixed-effects models were estimated using the Satterthwaite approximation via the `lmerTest` package (Kuznetsova, Bruun Brockhoff, & Haubo Bojesen Christensen, 2016). Random effects (intercepts and slopes) for subjects and for items were included wherever appropriate, and are described below.

Data availability. Our data along with all methods, materials, and analysis scripts, are available in public repositories described on the Open Science Framework page for this research here: osf.io/3navm.

Results

Acoustic similarity increased through iteration. Imitations of environmental sounds became more stable over the course of being repeated as revealed by increasing acoustic similarity judgments along individual transmission chains. Acoustic similarity ratings were fit with a linear mixed-effects model predicting perceived acoustic similarity from generation with random effects (intercepts and slopes) for raters. To test whether the hypothesized increase in acoustic similarity was true across all seed sounds and categories, we added random effects (intercepts and slopes) for seed sounds nested within categories. The results showed that, across raters and seeds, imitations from later generations were rated as sounding more similar to one another than imitations from earlier generations, $b = 0.10$ (SE = 0.03), $t(11.9) = 3.03$, $p = 0.011$ (Fig. 2). This result suggests that imitations became more stable (i.e., easier to imitate with high fidelity) with each generation of repetition.




Figure 2. Change in perception of acoustic similarity over generations of iterated imitation. Points depict mean acoustic similarity ratings for pairs of imitations in each category. The predictions of the linear mixed-effects model are shown with ± 1 SE. Acoustic similarity increased over generations, indicating that repetition made the vocalizations easier to imitate with high fidelity.

Acoustic similarity was highest within transmission chains. Increasing similarity along transmission chains could also reflect the continuous degradation of the signal due to repeated imitation, in which case we would expect acoustic similarity to increase both within as well as between transmission chains as a function of generation of imitation. To rule out this alternative explanation, we calculated MFCCs for pairs of sounds sampled from within and between different transmission chains from consecutive generations across categories. To analyze the results, we fit a linear model predicting normalized acoustic similarity scores (z-scores) from the generation of sounds. A hierarchical model was not appropriate for this analysis because the between-chain pairs of sounds were sampled from

different categories, preventing any random effects due to category or seed from being included in the model. We found that acoustic similarity increased within chains more than it increased between chains, $b = -0.07$ (SE = 0.03), $t(6674.0) = -2.13$, $p = 0.033$ (Fig. S2). This result supports the conclusion that transmission chains were stabilizing on divergent acoustic forms as opposed to all chains converging on similar forms through continuous degradation.

Later generation imitations were transcribed more consistently. An additional test of stabilization and word-likeness was to measure whether later generation imitations were transcribed more consistently than first generation imitations. We collected a total of 2163 transcriptions — approximately 20 transcriptions per sound. Of these, 179 transcriptions (8%) were removed because they contained English words. Some examples of the final transcriptions are presented in Table 1.

To measure the similarity among transcriptions, we calculated the orthographic distance between the most frequent transcription and all other transcriptions of a given imitation. The orthographic distance measure was a ratio based on longest contiguous matching subsequences between pairs of transcriptions. We then fit a hierarchical linear model predicting orthographic distance from the generation of the imitation (First generation, Last generation) with random effects (intercepts and slopes) for seed sound nested within category¹. The results showed that transcriptions of last generation imitations were more similar to one another than transcriptions of first generation imitations, $b = -0.12$ (SE = 0.03), $t(3.0) = -3.62$, $p = 0.035$ (Fig. 3). The same result is reached through alternative measures of orthographic distance, such as the percentage of exact transcription matches for each imitation, $b = 0.10$ (SE = 0.03), $t(90.0) = 2.84$, $p = 0.006$, and the length of the longest matching substring, $b = 0.98$ (SE = 0.24), $t(15.1) = 4.14$, $p < 0.001$ (Fig. S3).

¹Random effects for subject were not appropriate because the distance measure was derived from pairwise comparisons of transcriptions generated by different transcribers. As a result, the degrees of freedom for the significance tests for the parameters of this model reflect the Satterthwaite approximation based on the number of seed sounds (16) nested within categories (4), not the number of unique transcribers ($N=216$).

206 Differences between transcriptions of human vocalizations and transcriptions directly of
207 environmental sounds are presented in the Supplementary Materials (Fig. S5).




Figure 3. Orthographic agreement among transcriptions of first and last generation imitations. Points depict the mean orthographic distance between the most frequent transcription and all other transcriptions of a given imitation, with error bars denoting ± 1 SE of the hierarchical linear model predictions. Transcriptions of later generation imitations were more similar to one another than transcriptions of first generation imitations, suggesting that repeating imitations made them easier to transcribe into English orthography than direct imitations of environmental sounds.

208 **Discussion**

209 Repeating imitations of environmental sounds over generations of unique speakers was
210 sufficient to create more wordlike forms even without any instruction to do so. We defined
211 wordlike-ness in terms of acoustic stability and orthographic agreement. With additional
212 repetitions, the acoustic forms of the imitations became more similar to one another,

indicating they became easier to repeat with high fidelity. The possibility that this similarity was due to uniform degradation across all transmission chains was ruled out by algorithmic analyses of acoustic similarity within and between chains demonstrating that acoustic similarity increased within chains but not between them. Additionally, later generation imitations were transcribed more consistently into English orthography, further supporting our hypothesis that repeating imitations makes them more word-like.

The results of Experiment 1 demonstrate the ease with which iterated imitation gives rise to unique word forms. However, the results do not address how these emergent words relate to the original sounds that were being imitated. As the imitations became more word-like, were they stabilizing on arbitrary acoustic and orthographic forms, or did they maintain some resemblance to the environmental sounds that motivated them? The purpose of Experiment 2 was to assess the extent to which repeated imitations and their transcriptions maintained a resemblance to the original set of seed sounds.

Experiment 2: Resemblance of imitations to original seed sounds

To assess the resemblance of repeated imitations to the original seed sounds, we measured the ability of participants naïve to the design of the experiment to match imitations and their transcriptions back to their original sound source relative to other seed sounds from either the same category or from different categories (Fig. 4). We used match accuracies to answer two questions concerning the effect of iterated imitation on resemblance to the original seed sounds. First, we asked whether and for how many generations the imitations and their transcriptions could be matched back to the original sounds. Second, we asked whether repeated imitation resulted in a uniform degradation of the signal in each imitation, or if repeated imitation resulted in some kinds of information degrading more rapidly than others. Specifically, we tested the hypothesis that if imitations were becoming more word-like, then they should also be interpreted more categorically, and thus we predicted that the imitations might lose individuating information that identifies the specific

239 source of an imitation more rapidly than category information that identifies the general
240 category of environmental sound being imitated.




Figure 4. Three types of matching questions used to assess the resemblance between the imitation (and transcriptions of imitations) and the original seed sounds. For each question, participants listened an imitation (dashed circles) or read a transcription of one, and had to guess which of 4 sound choices (solid circles) they thought the person was trying to indicate. True seed questions contained the specific sound that generated the imitation as one of the choices (the correct response). The remaining sound choices were sampled from different categories. Category match questions replaced the original seed sound with another sound from the same category. Specific match questions pitted the actual seed against the other seeds within the same category.

241 **Methods**

242 **Matching imitations to seed sounds.** Participants ($N=751$) recruited from
243 Amazon Mechanical Turk were paid to listen to imitations, one at a time, and for each one,
244 choose one of four possible sounds they thought the person was trying to imitate. The task

was unspeeded and no feedback was provided. Participants completed 10 questions at a time.

All 365 imitations were tested in each of the three question types depicted in Fig. 4. These questions differed in the relationship between the imitation and the four seed sounds provided as the choices in the question. Question types (True seed, Category match, Specific match) were assigned between-subject. Participants in the True seed and Category match conditions were provided four seed sounds from different categories as choices in each question. Participants in the Specific match condition were provided four seed sounds from the same category.

Matching transcriptions to seed sounds. Participants ($N=468$) recruited from Amazon Mechanical Turk completed a modified version of the matching survey described above. Instead of listening to imitations, participants now read a word (a transcription of an imitation), which they were told was an invented word. They were instructed that the word was invented to describe one of the four presented sounds, and they had to guess which one. The distractors for all questions were between-category, i.e. true seed and category match. Specific match questions were omitted.

Of the unique transcriptions that were generated for each sound (imitations and seed sounds), only the top four most frequent transcriptions were used in the matching experiment. Participants who failed a catch trial ($N=6$) were excluded, leaving 461 participants in the final sample.

Results

Imitations retained category information more than individuating information. Response accuracies in matching imitations to seed sounds were fit by a generalized linear mixed-effects model predicting match accuracy as different from chance (25%) based on the type of question being answered (True seed, Category match, Specific match) and the generation of the imitation. Question types were contrast coded using Category match questions as the baseline condition in comparison to the other two question

types each containing the actual seed that generated the imitation as one of the choices. The model included random intercepts for participant², and random slopes and intercepts for seed sounds nested within categories.

Accuracy in matching imitations to seed sounds was above chance for all question types for the first generation of imitations, $b = 1.65$ (SE = 0.14) log-odds, odds = 0.50, $z = 11.58$, $p < 0.001$, and decreased steadily over generations, $b = -0.16$ (SE = 0.04) log-odds, $z = -3.72$, $p < 0.001$. We then tested whether this increase in difficulty was constant across the three types of questions or if some question types became more difficult than others. The results are shown in Fig. 5A. Performance decreased over generations more rapidly for questions requiring a within-category distinction than for between-category questions, $b = -0.08$ (SE = 0.03) log-odds, $z = -2.68$, $p = 0.007$, suggesting that between-category information was more resistant to loss through repeated imitation.

An alternative explanation for this result is that the within-category match questions are simply more difficult because the sounds provided as choices are more acoustically similar to one another than the between-category questions, and therefore, performance might be expected to drop off more rapidly with repeated imitation for these more difficult questions³. However, performance also decreased for the easiest type of question where the correct answer was the actual seed generating the imitation (True seed questions; see Fig. 4); the advantage of having the true seed among between-category distractors decreased over generations, $b = -0.07$ (SE = 0.02) log-odds, $z = -2.77$, $p = 0.006$. The observed increase in the “category advantage” (i.e., the advantage of having between-category distractors)

²Random slopes for generation were not appropriate in the by-subject random effects because data collection was batched by generation of imitation, and therefore each participant did not sample across the range of generations.

³We observed that performance on some Specific match questions dropped below chance for later generations indicating participants had an apparent aversion to the nominally correct answer. Additional analyses showed that participants were not converging on a single incorrect response. The reason for this pattern is at present unclear. Removing these trials from the analysis does not substantively change the conclusions.

combined with a decrease in the “true seed advantage” (the advantage of having the actual seed among the choices), shows that the changes induced by repeated imitation caused the imitations to lose some of properties that linked the earlier imitations to the specific sound that motivated them, while nevertheless preserving a more abstract category-based resemblance.

Transcriptions retained information about seed sources. We next report the results of matching the written transcriptions of the auditory sounds back to the original environmental sounds. Remarkably, participants were able to guess the correct meaning of a word that was transcribed from an imitation that had been repeated up to 8 times, $b = 0.83$ (SE = 0.13) log-odds, odds = -0.18, $z = 6.46$, $p < 0.001$ (Fig. 5B). This was true for True seed questions containing the actual seed generating the transcribed imitation, $b = 0.75$ (SE = 0.15) log-odds, $z = 4.87$, $p < 0.001$, and for Category match questions where participants had to associate transcriptions with a particular category of environmental sounds, $b = 1.02$ (SE = 0.16) log-odds, $z = 6.39$, $p < 0.001$. The effect of generation did not vary across these question types, $b = 0.05$ (SE = 0.10) log-odds, $z = 0.47$, $p = 0.638$. The results of matching “transcriptions” directly of the environmental sounds are shown in Fig. S5.

Discussion

Even after being repeated up to 8 times, imitations retained a resemblance to the environmental sound that motivated them, even after being transcribed into orthographic forms. For imitations, but not for transcriptions, this resemblance was stronger for the category of environmental sound than the actual seed sound, suggesting that through repetition, the imitations were becoming more categorical. This result supports the results of Experiment 1 in demonstrating another aspect of wordlike-ness achieved through repeated imitation: Words, in addition to being stable in acoustic and orthographic forms, are also categorical, denoting all members of a category equally as opposed to identifying individual category members. Repeating imitations of environmental sounds is sufficient to remove

318 some of the individuating characteristics of the imitation while retaining a category-based
319 resemblance.

320 The reason the same effect was not observed in matching accuracy for transcriptions is
321 unknown. One possible reason is that the process of transcribing a non-linguistic
322 vocalization into a written word encourages transcribers to emphasize individuating
323 information about the vocalization. However, the fact that transcriptions of imitations can
324 be matched back to other category members (Category match questions) suggests that
325 transcriptions are still carrying some category information. Another possible reason is that
326 by subsetting the most frequent transcriptions, we unintentionally excluded less frequent
327 transcriptions that were more diagnostic of category information.

328 Experiments 1 and 2 document a process of gradual change from an imitation of an
329 environmental sound to a more wordlike form. But do these emergent words function like
330 other words in the language? In Experiment 3, we test the suitability of words taken from
331 the beginning and end of transmission chains in serving as category labels in a category
332 learning task.

333 **Experiment 3: Suitability of created words as category labels**

334 One consequence of imitations becoming more word-like is that they may make for
335 better category labels. For example, an imitation from a later generation, by virtue of having
336 a more word-like form, may be easier to learn as a label for the category of sounds that
337 motivated it than an earlier imitation, which is more closely yoked to a particular
338 environmental sound. To the extent that repeating imitations abstracts away the
339 idiosyncrasies of a particular category member (Edmiston & Lupyan, 2015; Lupyan &
340 Thompson-Schill, 2012), it may also be easier to generalize to new category members. We
341 tested these predictions using a category learning task in which participants learned novel
342 labels as category labels of the seed environmental sounds. The novel labels were
343 transcriptions of either first or last generation imitations gathered in Experiment 1.

Methods

Selecting words to learn as category labels. Our transmission chain design and subsequent transcription procedure created 1814 unique words. From these, we sampled words transcribed from first and last generation imitations, as well as transcriptions of the original seed sounds. Our procedure for sampling transcriptions to use as category labels was as follows: First, we removed transcriptions that contained less than 3 unique characters and transcriptions that were over 10 characters long. Of the remaining transcriptions, a sample of 56 were selected that were approximately equally associated with the target category. To measure the association between each imitation and its target category (the category of the seed sound), we used the match accuracy scores reported in Experiment 2. The reason for using this measure of association strength as a control for selecting words to learn as category labels was to be able to select words that were initially equally associated with the target categories. Equating along this dimension allowed for a more focused test of differences between the words in terms of generalization to new category members. The final sample of transcriptions were selected using a bootstrapping procedure which involved selecting a desired mean (the average association strength for eligible transcriptions of last generation imitations) and sampling transcriptions from first generation imitations and from seed sounds until the match accuracy of those imitations matched the desired mean within 1 standard deviation.

Procedure. Participants ($N=67$) were University of Wisconsin undergraduates who received course credit for participation. Participants were randomly assigned four novel labels to learn for four categories of environmental sounds. Full instructions are provided in the Supplementary Materials. Participants were assigned between-subject to learn labels (transcriptions) of either first or last generation imitations. Some participants learned labels from transcriptions of seed sounds (Fig. S6). On each trial, participants heard one of the 16 seed sounds. After a 1s delay, participants saw a label (one of the transcribed imitations) and responded yes or no using a gamepad controller depending on whether the sound and

the word went together. Participants received accuracy feedback (a bell sound and a green checkmark if correct; a buzzing sound and a red “X” if incorrect). Four outlier participants were excluded from the final sample due to high error rates and slow RTs.

Participants categorized all 16 seed sounds over the course of the experiment, but they learned them in blocks of 4 sounds at a time. Within each block of 24 trials, participants heard the same four sounds and the same four words multiple times, with a 50% probability of the sound matching the word on any given trial. At the start of a new block of trials, participants heard four new sounds they had not heard before, and had to learn to associate these new sounds with the words they had learned in the previous blocks.

Results

Later generation transcriptions yielded more efficient responding.

Participants began by learning through trial-and-error to associate four written labels with four categories of environmental sounds. The small number of categories made this an easy task (mean accuracy after the first block of 24 trials was 81%; Fig. S4). Participants learning transcriptions of first or last generation imitations did not differ in overall accuracy, $p = 0.887$, or reaction time, $p = 0.616$. After this initial learning phase (i.e. after the first block of trials), accuracy performance quickly reached ceiling and did not differ between groups $p = 0.775$. However, the response times of participants learning last generation transcriptions declined more rapidly with practice than participants learning first generation transcriptions, $b = -114.13$ (SE = 52.06), $t(39.9) = -2.19$, $p = 0.034$ (Fig. 6A). These faster responses suggest that, in addition to becoming more stable both in terms of acoustic and orthographic properties, repeating imitations makes them easier to process as category labels. We predict that given a harder task (i.e., more than four categories and 16 exemplars) would yield differences in initial learning rates as well.

Later generation transcriptions were better generalized.

Next, we examined whether transcriptions from last generation imitations were easier to generalize to novel

category exemplars. To test this hypothesis, we compared RTs on trials immediately prior to the introduction of novel sounds (new category members) and the first trials after the block transition (± 6 trials). The results revealed a reliable interaction between the generation of the transcribed imitation and the block transition, $b = -110.77$ (SE = 52.84), $t(39.7) = -2.10$, $p = 0.042$ (Fig. 6B). This result suggests that transcriptions from later generation imitations were easier to generalize to new category members.

Discussion

The results of a simple category learning experiment demonstrate a possible benefit to the stabilization of repeated imitations on more wordlike forms. As a consequence of being more wordlike, repeated imitations were responded to more quickly, and generalized to new category members more easily. These results suggest an advantage to repeating imitations from the perspective of the language learner in that they afford better category generalization.

General Discussion

Imitative words are found across the spoken languages of the world (Dingemanse et al., 2015; Imai & Kita, 2014; Perniss et al., 2010). Counter to past assumptions about the limitations of human vocal imitation, people are surprisingly effective at using vocal imitation to represent and communicate about the sounds in their environment (Lemaitre et al., 2016) and more abstract meanings (Perlman et al., 2015), making the hypothesis that early spoken words originated from imitations a plausible one. We examined whether simply repeating an imitation of an environmental sound—with no intention to create a new word or even to communicate—produces more word-like forms.

Our results show that through simple repetition, imitative vocalizations became more word-like both in form and function. In form, the vocalizations gradually stabilized over generations, becoming more similar from imitation to imitation. They also became increasingly standardized in accordance with English orthography, as later generations were

WORDS FROM IMITATIONS

21

more consistently transcribed into English words, providing converging evidence of stabilization. In function, the increasingly word-like forms became more effective as category labels. In a category learning experiment, naïve participants were faster at matching category labels derived from later-generation imitations than those derived directly from imitations of environmental sounds. This fits with previous research showing that the relatively arbitrary forms that are typical of words (e.g. “dog”) makes them better suited to function as category labels compared to direct auditory cues (Boutonnet & Lupyan, 2015; Edmiston & Lupyan, 2015; e.g. the sound of a dog bark; Lupyan & Thompson-Schill, 2012).

Even as the vocalizations became more word-like, they nevertheless maintained an imitative quality. After eight generations they could no longer be matched to the particular sound from which they originated any more accurately than they could be matched to the general category of environmental sound. Thus, information that distinguished an imitation from other sound categories was more resilient to transmission decay than exemplar information within a category. Remarkably, even after the vocalizations were transcribed into English orthography, participants were able to guess their original sound category from the written “words”. In contrast to the vocalizations, participants continued to be more accurate at matching late generation transcriptions back to their particular source sound relative to other exemplars from the same category.

Although the number of imitative words in contemporary languages may appear to be very small (Crystal, 1987; Newmeyer, 1992), increasing evidence from disparate languages shows that vocal imitation is, in fact, a widespread source of vocabulary. Cross-linguistic surveys indicate that onomatopoeia—imitative words used to represent sounds—are a universal lexical category found across the world’s languages (Dingemanse, 2012). Even English, a language that has been characterized as relatively limited in iconic vocabulary (Vigliocco, Perniss, & Vinson, 2014), is documented as having hundreds of clearly imitative words including words for human and animal vocalizations as well as various types of environmental sounds (Rhodes, 1994; Sobkowiak, 1990). Besides words that are directly

imitative of sounds—the focus of the present study — many languages contain semantically broader inventories of ideophones. These words comprise a grammatically and phonologically distinct class of words that are used to express various sensory-rich meanings, such as qualities related to manner of motion, visual properties, textures and touch, inner feelings and cognitive states (Dingemanse, 2012; Nuckolls, 1999; Voeltz & Kilian-Hatz, 2001). As with onomatopoeia, ideophones are often recognized by naïve speakers as bearing a degree of resemblance to their meaning (Dingemanse, Schuerman, & Reinisch, 2016).

Our study focused on imitations of environmental sounds and more work remains to be done to determine the extent to which vocal imitation can ground de novo vocabulary creation in other semantic domains (Lupyan & Perlman, 2015; e.g., Perlman et al., 2015). What the present results make clear is that the transition from imitation to word can be a rapid and simple process: the mere act of iterated imitation can drive vocalizations to become more word-like in both form and function. Notably, just as onomatopoeia and ideophones of natural languages maintain a resemblance to the quality they represent, the present vocal imitations transitioned to words while retaining a resemblance to the original sound that motivated them.

References

- Arbib, M. A. (2012). *How the brain got language: The mirror system hypothesis* (Vol. 16). Oxford University Press.
- Armstrong, D. F., & Wilcox, S. (2007). *The gestural origin of language*. Oxford University Press.
- Bates, D., Mächler, M., Bolker, B., & Walker, S. (2015). Fitting Linear Mixed-Effects Models Using lme4. *Journal of Statistical Software*, 67(1), 1–48.
- Boutonnet, B., & Lupyan, G. (2015). Words Jump-Start Vision: A Label Advantage in Object Recognition. *Journal of Neuroscience*, 35(25), 9329–9335.
- Brown, R. W., Black, A. H., & Horowitz, A. E. (1955). Phonetic symbolism in natural

WORDS FROM IMITATIONS

23

languages. *Journal of Abnormal Psychology*, 50(3), 388–393.

Clark, H. H., & Gerrig, R. J. (1990). Quotations as demonstrations. *Language*, 66, 764–805.

Corballis, M. C. (2003). *From hand to mouth: The origins of language*. Princeton University Press.

Crystal, D. (1987). *The Cambridge Encyclopedia of Language* (Vol. 2). Cambridge Univ Press.

Dingemanse, M. (2012). Advances in the Cross-Linguistic Study of Ideophones. *Language and Linguistics Compass*, 6(10), 654–672.

Dingemanse, M. (2014). Making new ideophones in Siwu: Creative depiction in conversation. *Pragmatics and Society*.

Dingemanse, M., Blasi, D. E., Lupyan, G., Christiansen, M. H., & Monaghan, P. (2015). Arbitrariness, Iconicity, and Systematicity in Language. *Trends in Cognitive Sciences*, 19(10), 603–615.

Dingemanse, M., Schuerman, W., & Reinisch, E. (2016). What sound symbolism can and cannot do: Testing the iconicity of ideophones from five languages. *Language*, 92.

Donald, M. (2016). Key cognitive preconditions for the evolution of language. *Psychonomic Bulletin & Review*, 1–5.

Edmiston, P., & Lupyan, G. (2015). What makes words special? Words as unmotivated cues. *Cognition*, 143(C), 93–100.

Gamer, M., Lemon, J., Fellows, I., & Singh, P. (2012). *irr: Various Coefficients of Interrater Reliability and Agreement*.

Goldin-Meadow, S. (2016). What the hands can tell us about language emergence. *Psychonomic Bulletin & Review*, 24(1), 1–6.

Hall, K. C., Allen, B., Fry, M., Mackie, S., & McAuliffe, M. (2016). Phonological CorpusTools. *14th Conference for Laboratory Phonology*.

Hewes, G. W. (1973). Primate Communication and the Gestural Origin of Language.

WORDS FROM IMITATIONS

24

Current Anthropology, 14(1/2), 5–24.

Hockett, C. F. (1978). In search of Jove's brow. *American Speech*, 53(4), 243–313.

Imai, M., & Kita, S. (2014). The sound symbolism bootstrapping hypothesis for language acquisition and language evolution. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 369(1651).

Kendon, A. (2014). Semiotic diversity in utterance production and the concept of 'language'. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 369(1651), 20130293–20130293.

Klima, E. S., & Bellugi, U. (1980). *The signs of language*. Harvard University Press.

Kuznetsova, A., Bruun Brockhoff, P., & Haubo Bojesen Christensen, R. (2016). *lmerTest: Tests in Linear Mixed Effects Models*.

Lemaitre, G., & Rocchesso, D. (2014). On the effectiveness of vocal imitations and verbal descriptions of sounds. *The Journal of the Acoustical Society of America*, 135(2), 862–873.

Lemaitre, G., Houix, O., Voisin, F., Misdariis, N., & Susini, P. (2016). Vocal Imitations of Non-Vocal Sounds. *PloS One*, 11(12), e0168167–28.

Lewis, J. (2009). As well as words: Congo Pygmy hunting, mimicry, and play. In *The cradle of language*. The cradle of language.

Lupyan, G., & Perlman, M. (2015). The vocal iconicity challenge! In *The th biennial protolanguage conference*. Rome, Italy.

Lupyan, G., & Thompson-Schill, S. L. (2012). The evocative power of words: Activation of concepts by verbal and nonverbal means. *Journal of Experimental Psychology: General*, 141(1), 170–186.

Newmeyer, F. J. (1992). Iconicity and generative grammar. *Language*.

Nuckolls, J. B. (1999). The case for sound symbolism. *Annual Review of Anthropology*, 28(1), 225–252.

Perlman, M., Dale, R., & Lupyan, G. (2015). Iconicity can ground the creation of vocal

- 529 symbols. *Royal Society Open Science*, 2(8), 150152–16.
- 530 Perniss, P., Thompson, R. L., & Vigliocco, G. (2010). Iconicity as a General Property of
531 Language: Evidence from Spoken and Signed Languages. *Frontiers in Psychology*, 1.
532 Pinker, S., & Jackendoff, R. (2005). The faculty of language: what's special about it?
533 *Cognition*, 95(2), 201–236.
- 534 Rhodes, R. (1994). Aural images. *Sound Symbolism*, 276–292.
- 535 Shrout, P. E., & Fleiss, J. L. (1979). Intraclass correlations: uses in assessing rater reliability.
536 *Psychological Bulletin*, 86(2), 420–428.
- 537 Sobkowiak, W. (1990). On the phonostatistics of English onomatopoeia. *Studia Anglica*
538 *Posnaniensia*, 23, 15–30.
- 539 Tomasello, M. (2010). *Origins of human communication*. MIT press.
- 540 Vigliocco, G., Perniss, P., & Vinson, D. (2014). Language as a multimodal phenomenon:
541 implications for language learning, processing and evolution. *Philosophical*
542 *Transactions of the Royal Society B: Biological Sciences*, 369(1651),
543 20130292–20130292.
- 544 Voeltz, F. E., & Kilian-Hatz, C. (2001). *Ideophones* (Vol. 44). John Benjamins Publishing.

WORDS FROM IMITATIONS

26

Table 1

Examples of words transcribed from imitations.

Category	Seed	First generation	Last generation
glass	1	tingtingting	deetdedededeet
glass	2	chirck	correcto
glass	3	dirrng	wayew
glass	4	boonk	baroke
tear	1	scheeept	cheecheea
tear	2	feeshefee	cheeoooo
tear	3	hhhweerrr	chhhhhhewwwe
tear	4	ccccchhhhyeaahh	shhhhh
water	1	boococucuwich	eeverlusha
water	2	chwoochwooochwooo	cheiopshpshcheiopsh
water	3	atoadelchoo	mowah
water	4	awakawush	galonggalong
zipper	1	euah	izoo
zipper	2	zoop	veeeep
zipper	3	arrgt	owww
zipper	4	bzzzzup	izzip

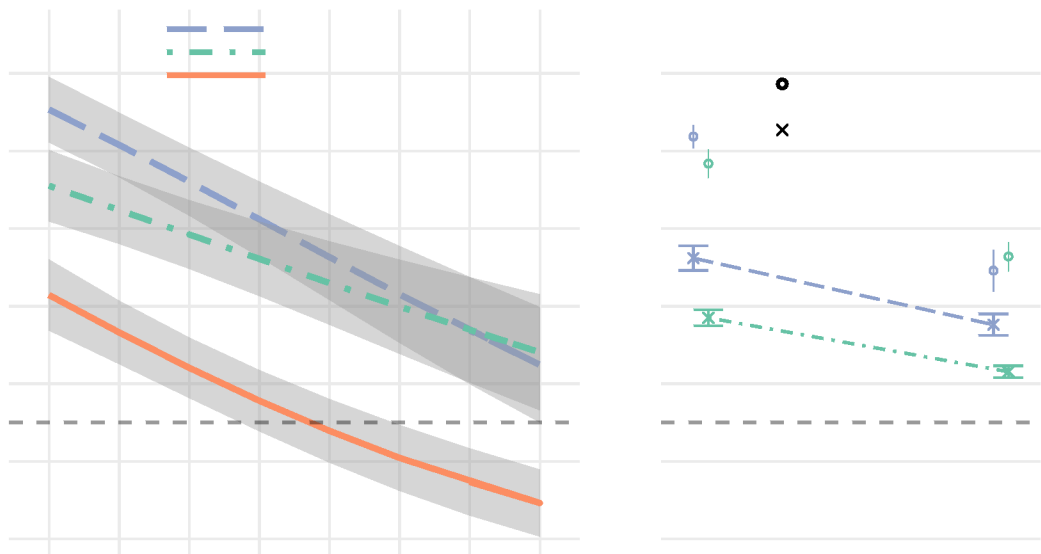


Figure 5. Repeated imitations retained category resemblance. A. Accuracy of matching vocal imitations to original seed sounds as a function of the generation during which the imitation was produced. Curves show predictions of the generalized linear mixed effects models with ± 1 SE of the model predictions. The “category advantage” (Category match vs. Specific match) increased over generations, while the “true seed advantage” (True seed v. Category match) decreased (see main text), suggesting that imitations lose within-category information more rapidly than between-category information. B. Accuracy of matching transcriptions of the imitations to original seed sounds (e.g., “boococucuwich” to a water splashing sound). Transcriptions of imitations could still be matched back to the category of sound that motivated the original imitation even after 8 generations. Circles show mean matching accuracy for the corresponding vocal imitations for comparison.




Figure 6. Repeated imitations made for better category labels. Participants learned novel labels (transcriptions of first or last generation imitations) for categories of environmental sounds. A. Mean RTs for correct responses in the category learning experiment with ± 1 SE. Participants achieved faster RTs in matching transcribed labels to environmental sounds for labels transcribed from later compared to earlier generation imitations. B. Cost of generalizing to new category members with ± 1 SE. After each block of trials, new environmental sounds were introduced, requiring participants to generalize the previously learned category labels to new category members. There was a generalization cost for the first generation labels, but not the last generation labels.

545

Table captions

546

Table 1. Examples of words transcribed from imitations.

Figure captions

Figure 1. Vocal imitations collected in the transmission chain experiment. Seed sounds (16) were sampled from four categories of environmental sounds: glass, tear, water, zipper. Participants imitated each seed sound, and then the next generation of participants imitated the imitations, and so on, for up to 8 generations. Chains are unbalanced due to random assignment and the exclusion of some low quality recordings.

Figure 2. Change in perception of acoustic similarity over generations of iterated imitation. Points depict mean acoustic similarity ratings for pairs of imitations in each category. The predictions of the linear mixed-effects model are shown with ± 1 SE. Acoustic similarity increased over generations, indicating that repetition made the vocalizations easier to imitate with high fidelity.

Figure 3. Orthographic agreement among transcriptions of first and last generation imitations. Points depict the mean orthographic distance between the most frequent transcription and all other transcriptions of a given imitation, with error bars denoting ± 1 SE of the hierarchical linear model predictions. Transcriptions of later generation imitations were more similar to one another than transcriptions of first generation imitations, suggesting that repeating imitations made them easier to transcribe into English orthography than direct imitations of environmental sounds.

568 *Figure 4.* Three types of matching questions used to assess the resemblance between
569 the imitation (and transcriptions of imitations) and the original seed
570 sounds. For each question, participants listened an imitation (dashed
571 circles) or read a transcription of one, and had to guess which of 4 sound
572 choices (solid circles) they thought the person was trying to indicate.
573 True seed questions contained the specific sound that generated the
574 imitation as one of the choices (the correct response). The remaining
575 sound choices were sampled from different categories. Category match
576 questions replaced the original seed sound with another sound from the
577 same category. Specific match questions pitted the actual seed against
578 the other seeds within the same category.

579 *Figure 5.* Repeated imitations retained category resemblance. A. Accuracy of
580 matching vocal imitations to original seed sounds as a function of the
581 generation during which the imitation was produced. Curves show pre-
582 dictions of the generalized linear mixed effects models with ± 1 SE of the
583 model predictions. The “category advantage” (Category match vs. Spe-
584 cific match) increased over generations, while the “true seed advantage”
585 (True seed v. Category match) decreased (see main text), suggesting that
586 imitations lose within-category information more rapidly than between-
587 category information. B. Accuracy of matching transcriptions of the
588 imitations to original seed sounds (e.g., “boococucuwich” to a water
589 splashing sound). Transcriptions of imitations could still be matched
590 back to the category of sound that motivated the original imitation
591 even after 8 generations. Circles show mean matching accuracy for the
592 corresponding vocal imitations for comparison.

Figure 6. Repeated imitations made for better category labels. Participants learned novel labels (transcriptions of first or last generation imitations) for categories of environmental sounds. A. Mean RTs for correct responses in the category learning experiment with ± 1 SE. Participants achieved faster RTs in matching transcribed labels to environmental sounds for labels transcribed from later compared to earlier generation imitations. B. Cost of generalizing to new category members with ± 1 SE. After each block of trials, new environmental sounds were introduced, requiring participants to generalize the previously learned category labels to new category members. There was a generalization cost for the first generation labels, but not the last generation labels.